

Activation Functions

Functions for introducing non-linearity in models

Input \rightarrow Activation \rightarrow Output

Why many activation functions?

- ▶ **Non-linearity**: Different ways of introducing non-linearity for learning complex patterns
- ▶ **Gradient flow**: Different backpropagation dynamics

@AlinMinutes



Linear Activation Function

The identity function

$$f(x) = x$$

Terms:

- ▶ **Linear:** Output proportional to input
- ▶ **Identity:** Returns same value as input


Key Properties:

- ▶ No non-linearity → Cannot learn complex patterns
- ▶ Network behaves as single linear transformation
- ▶ Used in output layer for regression problems



@AlinMinutes

Linear Activation Function



@AlinMinutes



Heaviside Step Function

Binary activation for early perceptrons

$$H(x) = \begin{cases} 0, & x < 0 \\ 1, & x \geq 0 \end{cases}$$

Terms:

- ▶ **Step Function:** Jumps from 0 to 1 at $x = 0$
- ▶ **Binary Activation:** Makes yes/no decisions

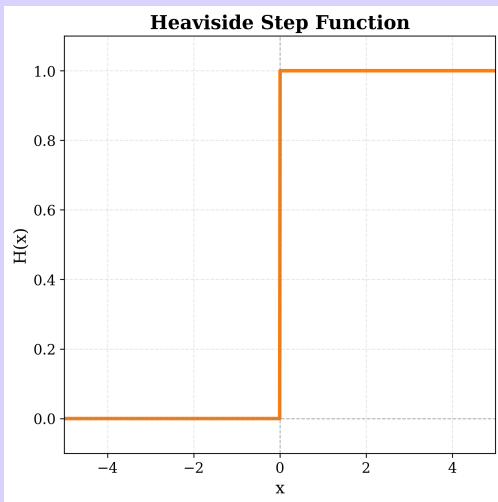
Key Properties:

- ▶ Used in McCulloch-Pitts neurons and perceptrons
- ▶ Not differentiable at $x = 0$
- ▶ Incompatible with gradient-based learning



@AlinMinutes

Heaviside Step Function



@AlinMinutes



Sigmoid Activation Function

Smooth S-shaped squashing function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Terms:

- ▶ **Sigmoid Curve:** S-shaped, maps to (0,1)
- ▶ **Derivative:** $\sigma'(x) = \sigma(x)(1 - \sigma(x))$

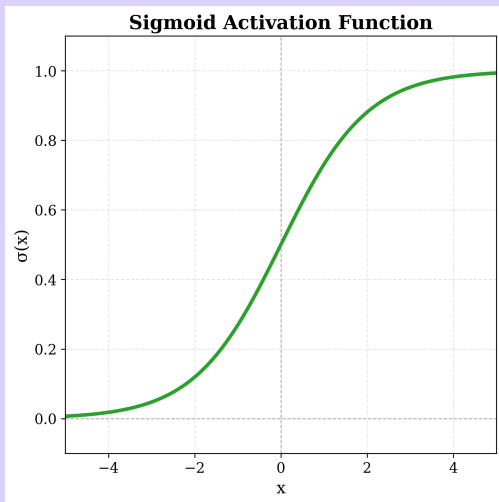
Key Properties:

- ▶ Used in logistic regression and early networks
- ▶ Probability interpretation for classification
- ▶ Suffers from vanishing gradient problem



@AlinMinutes

Sigmoid Activation Function



@AlinMinutes



Tanh Activation Function

Zero-centered sigmoid variant

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

Terms:

- ▶ Hyperbolic Tangent: Maps to $(-1,1)$
- ▶ Derivative: $\tanh'(x) = 1 - \tanh^2(x)$


Key Properties:

- ▶ Zero-centered outputs improve convergence
- ▶ Used in LSTMs for candidate memory updates
- ▶ Still suffers from gradient saturation



@AlinMinutes

Tanh Activation Function



@AlinMinutes



ReLU (Rectified Linear Unit)

Simple, efficient non-linearity

$$f(x) = \max(0, x)$$

Terms:


- ▶ **Rectified:** Zeroes out negative values
- ▶ **Derivative:** $f'(x) = \begin{cases} 1, & x > 0 \\ 0, & x < 0 \\ \text{undefined}, & x = 0 \end{cases}$

Key Properties:

- ▶ Computationally efficient
- ▶ Helps mitigate vanishing gradients
- ▶ Suffers from "dying ReLU" problem



ReLU Activation Function



@AlinMinutes



Leaky ReLU

ReLU with small negative slope

$$f(x) = \begin{cases} x, & x > 0 \\ \alpha x, & x \leq 0 \end{cases}$$

Terms:

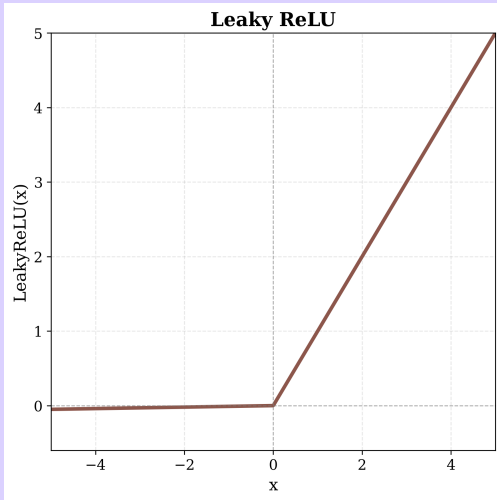
- ▶ Leak Factor α : Small constant (e.g., 0.01)
- ▶ Modified ReLU: Allows small negative outputs

Key Properties:

- ▶ Prevents "dying ReLU" problem
- ▶ Maintains gradient flow for negative inputs



Leaky ReLU Activation Function



@AlinMinutes

Parametric ReLU (PReLU)

ReLU with learnable negative slope(s)

$$f(x) = \begin{cases} x, & \text{if } x > 0 \\ \alpha x, & \text{if } x \leq 0 \end{cases}$$

Terms:

- ▶ α (Learnable Slope(s)): Trained parameter(s)
- ▶ Neuron/Layer/Global α : Different granularities

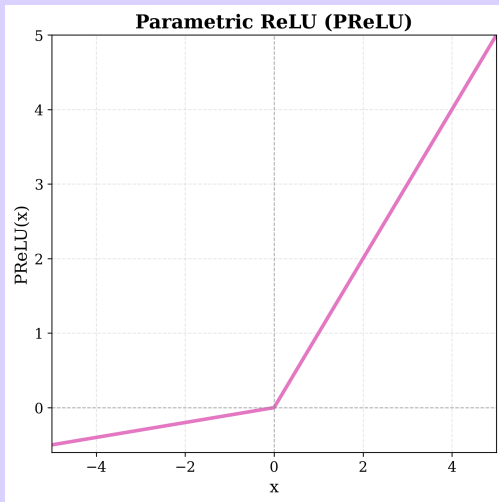
Key Properties:

- ▶ Adapts negative slopes to data patterns
- ▶ More flexible than standard ReLU
- ▶ Prevents "dying ReLU" problem



@AlinMinutes

Parametric ReLU Activation Function



@AlinMinutes



Exponential Linear Unit (ELU)

Smooth activation with saturation for negatives

$$f(x) = \begin{cases} x, & \text{if } x > 0 \\ \alpha(e^x - 1), & \text{if } x \leq 0 \end{cases}$$

Terms:


- ▶ α : Controls curve for negative inputs
- ▶ **Exponential Component**: Gradual negative values

Key Properties:

- ▶ Smooth and differentiable everywhere
- ▶ Reduces bias shift during training
- ▶ Negative values saturate to $-\alpha$



ELU Activation Function



@AlinMinutes



Swish- β

Non-monotonic with smooth gradients

$$f(x) = x \cdot \sigma(\beta x) = x \cdot \frac{1}{1 + e^{-\beta x}}$$

Terms:

- ▶ $\sigma(x)$: Sigmoid function
- ▶ β : Learnable parameter controlling non-linearity

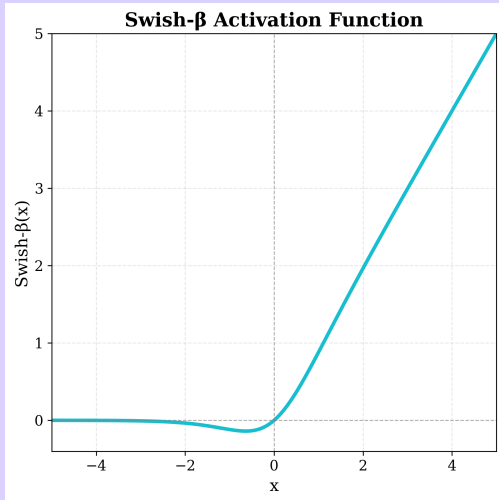
Key Properties:

- ▶ When $\beta = 1$, becomes standard Swish/SiLU
- ▶ Smooth and differentiable everywhere
- ▶ Handles negative values (non-monotonically) for better gradient flow (controlled by β)



@AlinMinutes

Swish- Activation Function



@AlinMinutes



SiLU (Sigmoid Linear Unit)

Also known as Swish-1 activation

$$f(x) = x \cdot \sigma(x) = \frac{x}{1 + e^{-x}}$$

Terms:

- ▶ $\sigma(x)$: Standard sigmoid function
- ▶ **Swish-1**: Special case of Swish where $\beta = 1$

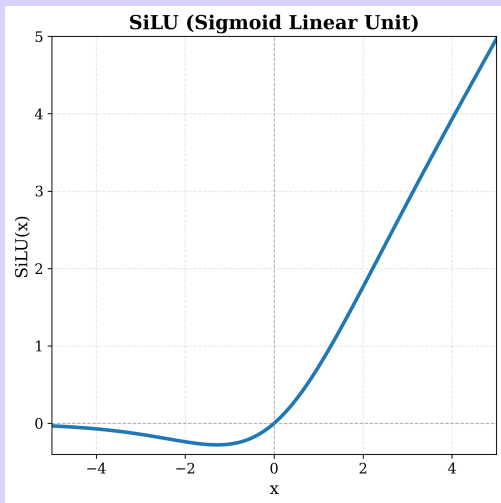
Key Properties:

- ▶ Smooth, non-monotonic activation function
- ▶ Unbounded above, bounded below
- ▶ Used in modern neural networks



@AlinMinutes

SiLU Activation Function



@AlinMinutes



GeLU (Gaussian Error Linear Unit)

Uses normal CDF for modulating inputs

$$f(x) = x \cdot \Phi(x) \approx x \cdot \frac{1}{2} \left(1 + \operatorname{erf} \left(\frac{x}{\sqrt{2}} \right) \right)$$

Terms:


- ▶ $\Phi(x)$: CDF of standard normal distribution
- ▶ erf : Error function related to normal distribution

Key Properties:

- ▶ Smooth and differentiable
- ▶ Probabilistic interpretation of activation
- ▶ Used in transformer architectures (e.g., GPT)



GeLU Activation Function



@AlinMinutes



GLU (Gated Linear Unit)

Input controls its own flow

$$\text{GLU}(x) = \sigma(xW + b) \odot (xV + c)$$

Terms:

- ▶ W, V : Learnable weight matrices
- ▶ b, c : Bias vectors
- ▶ $\sigma(x)$: Sigmoid activation function

Key Properties:

- ▶ Selective information flow via gating
- ▶ Smooth gradient flow with no dead zones
- ▶ Widely used in language models



@AlinMinutes

SwiGLU

Advanced gating using Swish

$$\text{SwiGLU}(X) = \text{Swish}_{\beta}(XW_1) \odot (XW_2)$$

Terms:

- ▶ $\text{Swish}_{\beta}(x) = x \cdot \sigma(\beta x)$
- ▶ W_1, W_2 : Learnable weight matrices
- ▶ \odot : Element-wise multiplication

Key Properties:

- ▶ Smoother gating than standard GLU
- ▶ Better information flow and gradient propagation
- ▶ Used in modern transformers (LLaMA, PaLM)



@AlinMinutes